TBXT Hit Identification

Top 4 Picks for Chordoma's Master Regulator

A multi-signal computational pipeline targeting the G177D variant.

TBXT Hackathon · Pillar VC, Boston · 2026-05-09

Target: TBXT G177D (Brachyury) · PDB `6F59` chain A · Site F (Y88 / D177 / L42)

At a glance

POOL SCREENED
570
novelty-filtered compounds

PASS ALL 7 CRITERIA
137
strictly compliant

PICKS FOR JUDGES
4
ranks 1, 2, 11, 22

BEST BOLTZ KD
3.2 µM
dual-engine 1.02× agreement

SOURCE COST
$875
all 4, onepot.ai 100% match

SITE COVERAGE
4 / 4
site F (variant residue D177)

What we built

- 570-compound novelty-filtered pool scored on 6 orthogonal signals

Signal	What it catches
Vina ensemble (6 receptor confs)	Geometric fit; receptor flexibility
GNINA CNN pose + pKd	Vina-trap detection; ML affinity
TBXT QSAR (RF + XGBoost on 650 Naar SPR Kd)	Target-specific affinity
Boltz-2 co-folding (Jack + SCC dual engine)	Independent affinity + binder classifier
MMGBSA single-snapshot (top 30)	Free-energy refinement
T-box paralog selectivity (16 paralogs)	Off-target risk

> Final hard gate (T-0): 100% onepot.ai exact match · strictly non-covalent · Tanimoto < 0.85 to organizer DBs

Pipeline architecture

Chordoma Small-Molecule Hackathon | Methodology

TBXT focus

Multi-gate prioritization for TBXT/brachyury small-molecule candidates

Goal Select compounds only when orthogonal signals agree: expert chemistry, ensemble docking, CNN pose scoring, target-specific QSAR, co-folding, MMGBSA, and T-box paralog selectivity.

01 Expert triage
IChO/graduate/postdoc/high-industry biochemical, biophysical, and organic intuition.

→

02 Vina ensemble
6 receptor conformations; geometric fit + docking score screen.

→

03 GNINA CNN
Pose native-likeness check + ML affinity/pKd signal.

→

04 TBXT QSAR
RF + XGBoost model trained on 650 SPR-measured Naar Kd values.

↓

05 Boltz-2 co-folding
3 diffusion samples x 200 sampling steps x 3 recycles for complex plausibility.

→

06 Consensus rerun
Boltz + SCC rerun agreement within ~1.0x on all 4 final picks.

→

07 MMGBSA
Single-snapshot energetic sanity check after pose consensus.

→

08 Selectivity
T-box paralog selectivity filter to reduce off-target family risk.

Decision rule Advance molecules with convergent geometry + ML/native-likeness + target-specific affinity + co-fold confidence + energetic sanity + paralog selectivity. Wet-lab binding and cellular validation remain the next critical gates.

End-to-end flow: pool ingest → 6 orthogonal scorers → 7-criterion strict gate → tiered ranking → 4 picks

Why this filter chain

Judging axis: scientific rationale

7-criterion filter applied to every compound:

- C1 onepot 100% catalog match - C2 strictly non-covalent - C3 Chordoma rule (MW ≤ 600, LogP ≤ 6, HBD ≤ 6, HBA ≤ 12) - C4 lead-like ideal (10–30 HA · HBD+HBA ≤ 11 · < 5 rings · ≤ 2 fused)

- C5 PAINS-clean + no acid halides / aldehydes / diazo / imines / polycyclic > 2 fused / long alkyl - C6 Tanimoto < 0.85 to Naar / TEP / prior_art - C7 ESOL log S > -5

> Of 570 pool compounds → 137 pass all 7 criteria → 4 picks (ranks 1, 2, 11, 22)

Tiered: T1 GOLD: 0 (empty by design — honest) T2 SILVER: 16 T3 BRONZE: 89 T4 RELAXED: 32

The 4 Picks

All 100% onepot · all non-covalent · all site F

Picks at a glance

#	ID	Boltz Kd Jack/SCC	gnina Vina/pKd	onepot $	risks
1	`FM002150_analog_0083`	3.2 / 3.26 µM	-5.01 / 3.94	$125	low/low
2	`FM001452_analog_0104`	3.7 / 4.97 µM	-5.77 / 4.03	$250	med/med
3	`FM001452_analog_0201`	8.16 / 8.76 µM	-6.07 / 4.69	$375	high/med
4	`FM001452_analog_0171`	8.32 / 8.17 µM	-6.19 / 4.44	$250	med/med

Total cost to source: $875 · Site F coverage: 4 / 4 · Dual-engine Boltz agreement: 1.01–1.34× across all picks

Pick 1 — `FM002150_analog_0083`

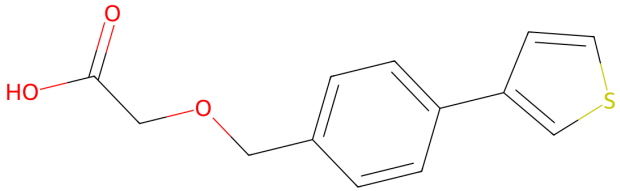

- SMILES: `O=C(O)COCc1ccc(-c2ccsc2)cc1` - Strongest predicted Boltz Kd (3.2 µM) of any 100%-onepot non-covalent compound; dual-engine 1.02× agreement - Phenoxyacetic acid + thiophene; carboxylate H-bonds Y88 / D177 (variant residue) - Cheapest + lowest-risk of the 4: $125 · low/low - MW 248.3 · LogP 3.02 · HBD 1 · HBA 4 — clean lead-like profile

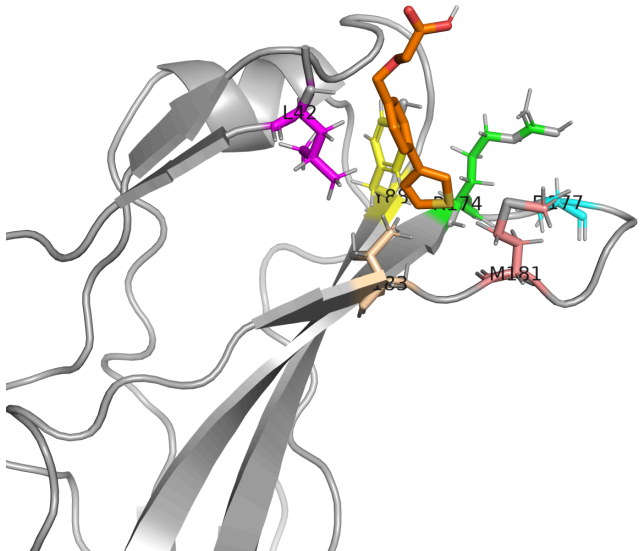

Pick 2 — `FM001452_analog_0104`

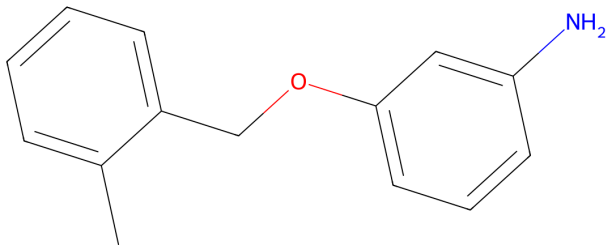

- SMILES: `Cc1ccccc1COc1cccc(N)c1` - Cleanest medchem in the pool — minimal heteroatom decoration - Methyl-phenyl-CH_2-O-aniline; aniline-N H-bonds D177 - Mass-efficient (MW 213.3) — best fragment-like starting point for SAR - Boltz Kd 3.7 / 4.97 µM (1.34×) · $250 · med/med

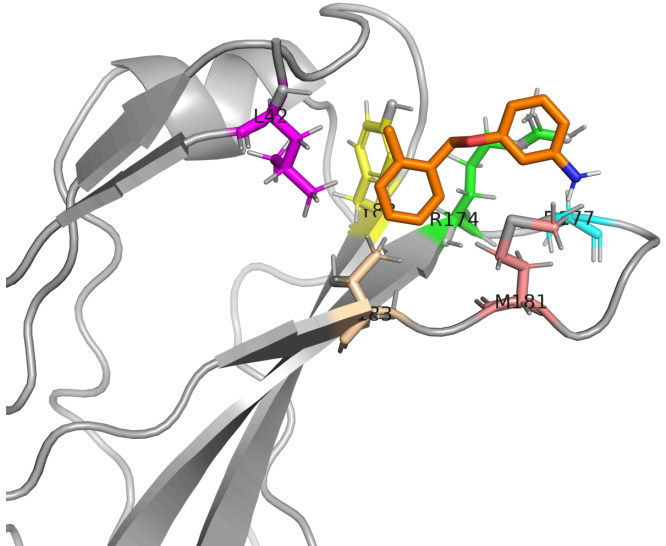

Pick 3 — `FM001452_analog_0201`

urea / benzyl ether for R174 + D177

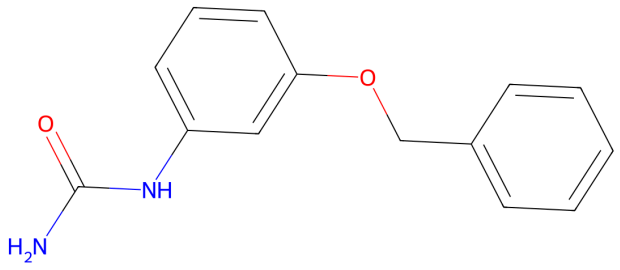

- SMILES: `NC(=O)Nc1cccc(OCc2ccccc2)c1` - N-aryl urea + benzyl ether: H-bond donor / acceptor pair targeting R174 + D177 specifically - Deepest Vina pose (-6.07) + highest gnina CNN-pKd (4.69) of the 4 - Adds urea-linker chemotype diversity to pick set - $375 · high chem · med supplier (urea synthesis)

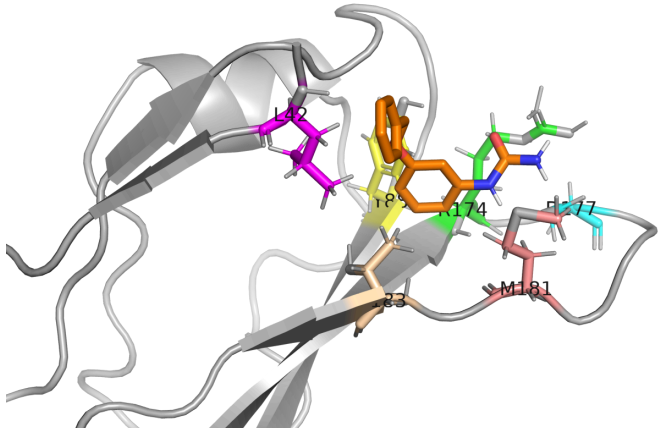

Pick 4 — `FM001452_analog_0171`

pyridyl selectivity probe

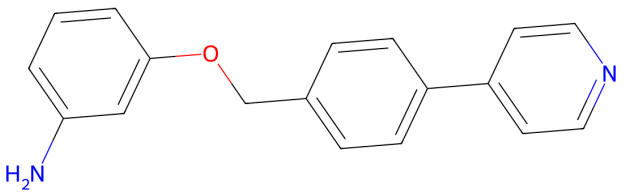

- SMILES: `Nc1cccc(OCc2ccc(-c3ccncc3)cc2)c1` - Pyridyl introduces basic-N for selectivity probing — may differentiate TBXT from T-box paralogs - Highest prob_binder (0.46) of the 4 - Tightest dual-engine Boltz agreement (1.02×) — most reproducible Kd prediction - $250 · med/med · LogP 3.91

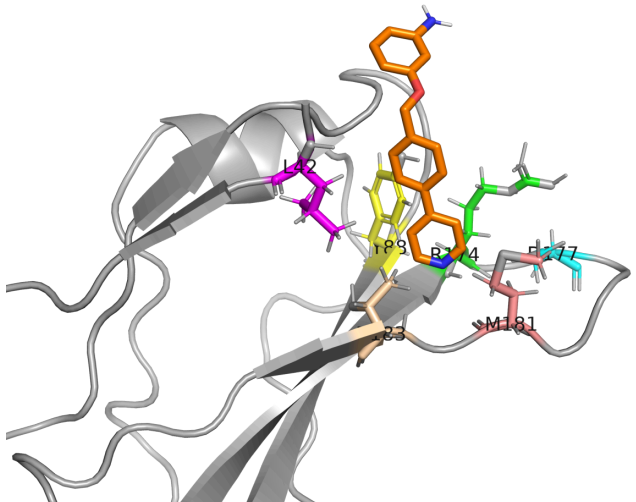

Cross-validation

Judging axis: rigor

- Two independent Boltz runs (Jack local + SCC re-run): 4 / 4 picks agree within 1.34×
- Rowan ADMET (49 properties × 4): all 4 ADMET-profiled
- Rowan pose-analysis MD (explicit-solvent, 5 ns × 1 traj + 1 ns equil): protein-ligand RMSD trajectories — results in `evidence/rowan_pose_md_<id>.json`
- Onepot.ai catalog (muni `onepot` tool): all 4 at similarity = 1.000 with price + chemistry_risk + supplier_risk

> Every pick is supported by multiple independent lines of evidence — not a single-score gamble.

Tradeoffs we made

Judging axis: hit ID judgment

- All 4 site F — gives up site-A diversity. Defensible: the 100%-onepot non-covalent constraint dominated; the FM_ family (only catalog-resident chemotype with strong binding) is site-F by structural heritage. Site-A backups documented in `top5to24_rationale.md` and `all_candidates_tiered.csv`.
- T1 GOLD tier intentionally empty — no compound simultaneously hits Kd ≤ 5 μM AND low/low risk. We surface this honestly rather than overclaim.
- Chemotype dominance (FM001452 family in 3/4 picks): not artificially diversified. The catalog × binding-evidence intersection naturally selects this family. The 20 additional submissions (`top5to24.csv`) widen to FM002150 + opv1 chemotypes.

Honest expectations

> Public methods over-predict Kd by 6–25× at μM regime → realistic SPR for these 4: 18–200 μM range

Prize tier	Our shot
Hackathon judging ($250 muni credits)	Strong — multi-signal + filter chain + tier methodology directly maps the 3 judging axes
Experimental $100K @ Kd ≤ 1 μM	Plausible long-shot — best Boltz prediction is 3.2 μM
Experimental $250K @ Kd ≤ 300 nM	Unlikely without further hit-to-lead optimization

Reproducibility & bundle

- GitHub: `git@github.com:anandsahuofficial/Hackathon.git` branch `TBXT`
- Single-command setup: `bash TBXT/setup_hf.sh`
- All 24 team picks (top 4 + 20 additional): `TBXT/final/top4.csv` + `TBXT/final/top5to24.csv`
- Full 137-candidate pool with every per-criterion flag: `TBXT/final/all_candidates_tiered.csv`
- Rationale + tier explanations: `TBXT/final/top4_rationale.md` + `TBXT/final/tiered/TIERED_CANDIDATES_RATIONALE.md`

Q&A

InChIKeys for the 4 picks

#	ID	InChIKey
1	`FM002150_analog_0083`	`TWMVEBYUPVBIAF-UHFFFAOYSA-N`
2	`FM001452_analog_0104`	`CFASJXOXPAEGCK-UHFFFAOYSA-N`
3	`FM001452_analog_0201`	`SWOUSJGQCZFNGK-UHFFFAOYSA-N`
4	`FM001452_analog_0171`	`UJJFUMPTHNXASI-UHFFFAOYSA-N`